

# Your Postgres DB performance Radar: **pgBadger**

Priyanka Chatterjee  
Schwarz IT KG - STACKIT

# Agenda

- 01** What is Pgbadger
- 02** Installation and execution
- 03** Options
- 04** Overview of a sample report
- 05** Use cases
- 06** Incremental and Cumulative reports
- 07** Nuiances

- pgBadger is free software distributed under the PostgreSQL Licence
- Utility to parse and analyze Postgres logs and create easy to understand reports
- Written in Perl and uses JavaScript library (flotr2) to create graphs
- It can parse huge log files as well as compressed files(gzip, bzip2, lz4, xz, zip and zstd)
- Target users: DBA, developer, anyone who uses Postgres
- Monitor DB performance, optimise queries, troubleshoot bottlenecks

- Supported on Multiple OS: Linux, MacOS, Windows
- Package manager

Eg: (ubuntu): # apt update

# apt install pgbadger

Github repo: <https://github.com/darold/pgbadger>

- **Postgres parameters required:**
- logging\_collector = on
- log\_min\_duration\_statement = 0 (or a value of choice)
- log\_checkpoints = on
- log\_connections = on
- log\_disconnections = on
- log\_lock\_waits = on
- log\_temp\_files = on
- log\_autovacuum\_min\_duration = 0
- log\_error\_verbosity = default
- log\_line\_prefix (must have at least specifies the %t (time) and %p (process))

- Creating a report: pgbadger postgresql.log
- Default output file: generated in current directory as out.html

Available Options: 104

- -b | --begin datetime
- -e | --end datetime
- -d | --dbname database
- -j | --jobs number
- -J | --Jobs number
- -m | --maxlength size
- -o | --outfile filename
- -O | --outdir path

# Overview (example report)

# Use case 1: query slowness

Where to look for in the report?

- Slowest individual query
- Most frequent query

# Use case 2: Temp files

Where to look for in the report

- Overview -> Global stats -> Temporary Files
- Queries generating the most temporary Files



# Use case 3: connections

Where to look for in the report

- Overview -> Global stats -> connections
- Connections per database
- Connections per user
- Events : Fatal like “Remaining connection slots are reserved for superuser connections” and/or “Sorry, too many clients already”

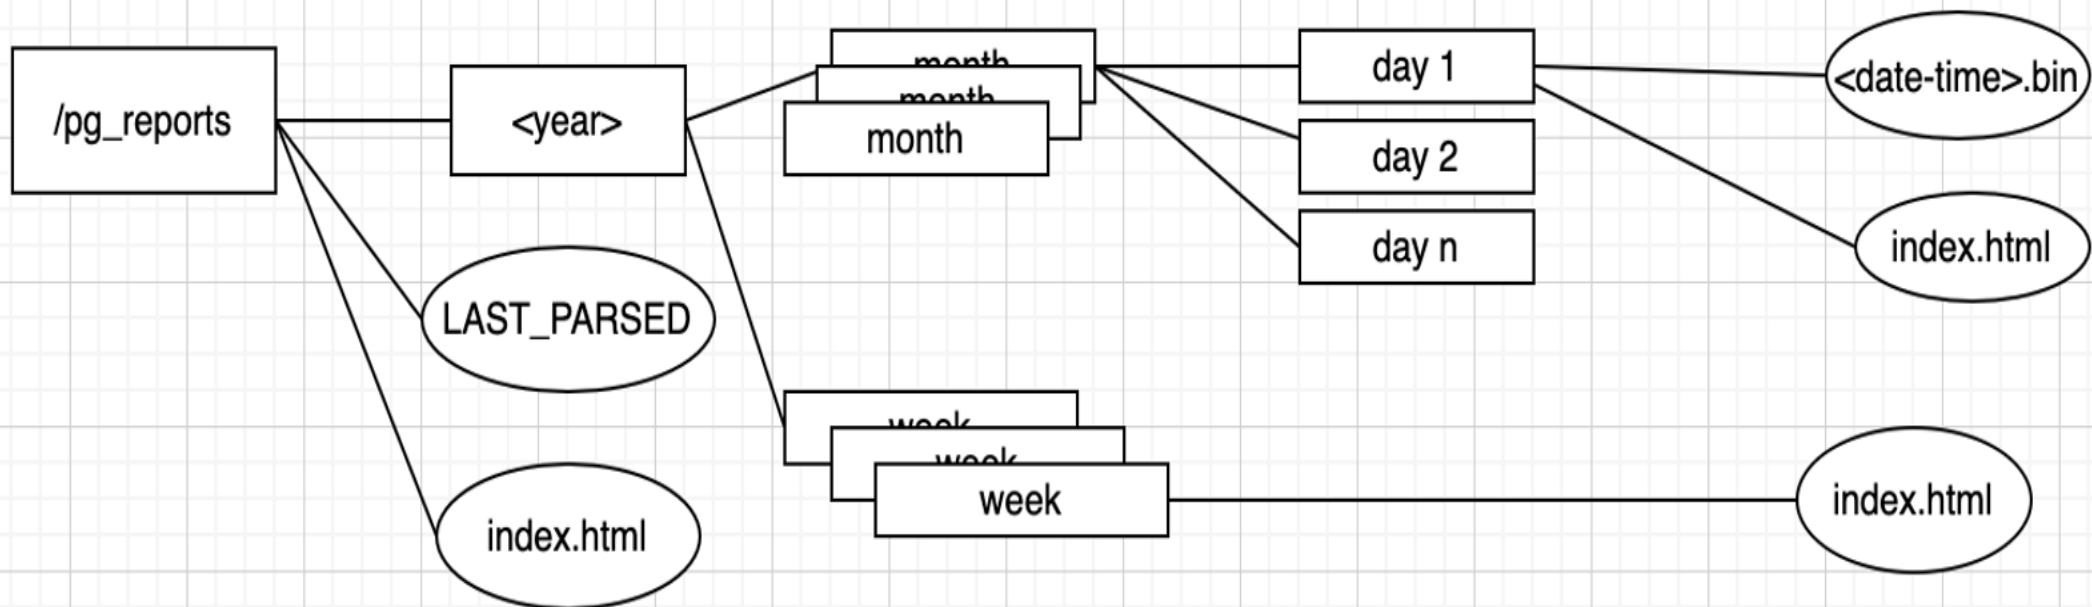
# Use case 4 : deadlocks

Where to look for in the report

- Events

# Incremental and Cumulative Reports

- -O | --outdir path
- -I | --incremental
- 0 1 \* \* \* \* /usr/bin/pgbadger -I -q /pg\_log/postgresql-\*.csv -O /pgbadger\_reports
- pgbadger --month-report 2024-11 /pgbadger\_reports



# Nuiances

- You need to install perl packages
- The graphs on the report are in accordance to your web browser timezone : -Z | --timezone +/-XX option
- If parallel processing, lot of temp files(tmp\_pgbadgerXXXX.bin) are generated in /tmp and cleaned up only on completion